

Ashish Reddy Jaddu
Full-Stack & Generative AI Engineer
Canada • ashishjaddu@gmail.com • +1 (514) 577-1298 • [LinkedIn](#) • [GitHub](#)

SUMMARY

Full-stack & AI engineer with startup CTO experience building LLM-powered legal research and workflow automation products for law firms. Led multiple product pivots from GPT-based legal research to sanction workflow automation, architecting secure Azure/GCP backends and building Next.js/React UIs that turn complex sanction and compliance workflows into software.

PROFESSIONAL EXPERIENCE

LEGISIMPLE

Chief Technology Officer

MONTREAL, CANADA

JULY 2024 – PRESENT

AI-native legal-tech startup building research and workflow automation tools for law firms.

- Own the technical roadmap and architecture, deciding on build-vs-buy and LLM vendors and working with founders and law-firm stakeholders to keep product and infrastructure aligned with legal workflows.
- Led four product pivots (**GPT-based legal research → Canadian case law search → US case law search → workflow automation for law firms**), shipping MVPs at each stage and validating them directly with lawyers.
- Built RAG-based legal research platform using NLP techniques for document parsing and GPT-powered semantic search across case law search MVPs for Canada and the US, using embeddings/vector search and Vertex AI fine-tuning (1,000 → 5,000 records, 82% → 87% accuracy) plus automated retraining pipelines driven by user feedback.
- Optimized Python/Django APIs with caching and query tuning (30% faster responses) and refactored backend services to reduce memory usage by 25%, while adding rate-limiting and Cloudflare protection to block abusive scraping and protect sensitive legal data.
- Built an end-to-end, multi-channel workflow automation platform for a partner law firm's sanction-progression process, where mobile apps capture complaints and deliver final documents while dashboard users review AI-driven analysis and record sanctions, cutting a 3-hour manual process to 15–30 minutes and logging all actions for audits and future queries.
- Architected an Azure workflow backend using Python (App Services, Durable Functions, Azure SQL) and Azure OpenAI + AI Search in a private VNet to secure law-firm data.
- Built and maintained Next.js/React frontends for legal research and sanction workflows, using Zustand for predictable client-side state and Azure SignalR for real-time dashboard updates so lawyers could see case and workflow progress live as actions were taken.
- Implemented Redis-based caching for expensive backend queries and instrumented the UI with Mixpanel and PostHog to track key events and drop-offs, improving perceived responsiveness and using analytics to refine critical workflow steps and reduce user friction.

HCL TECHNOLOGIES

Software Engineer

HYDERABAD, INDIA

OCT 2020 – JAN 2022

- **Developed and deployed RESTful APIs** powering internal React applications and **optimized MySQL queries**, improving page-load and interaction performance by ~20% and reducing DB response times by ~10%.
- **Debugged complex production issues** across frontend and backend services and **created a database-backed defect catalog**, streamlining issue tracking and boosting overall stability and QA effectiveness by ~40%.
- **Streamlined deployments for 5 web apps on Heroku**, improving configuration and resource usage to **cut hosting costs by ~15%** and make releases more reliable.

EDUCATION

CONCORDIA UNIVERSITY

Master of Science in Applied Computer Science

MONTREAL, CANADA

SEPT 2022 – APR 2024

TECHNICAL SKILLS

- **Programming Languages:** Python, Java, JavaScript (ES6+), TypeScript, SQL, PHP, C.
- **Web & Mobile:** Next.js, React.js, Node.js, Express.js, Django, HTML, CSS, React Native (prototyping), SwiftUI (prototyping), Zustand (state management), component-based UI, responsive UIs, REST API integration
- **Cloud & DevOps:** Azure (App Services, Durable Functions, Azure SQL, VNet, Azure OpenAI, Azure AI Search, Azure SignalR, Azure Translate), Google Cloud (Compute Engine, Cloud Run), Vercel, GitHub Actions, Docker, Kubernetes, monitoring & logging for production services.
- **Generative AI & LLMs:** OpenAI API (GPT-4, GPT-3.5-turbo), Azure OpenAI, Vertex AI, Groq API, LangChain, LangGraph, Retrieval-Augmented Generation (RAG), vector embeddings, semantic search, vector databases (Pinecone, Azure AI Search), prompt engineering, AI agents, LLM fine-tuning (API-based), natural language processing (NLP), production LLM integration, LLM evaluation and monitoring.
- **Data & Storage:** PostgreSQL, MySQL, MongoDB, SQLite, Neo4j, Pinecone, SQL performance tuning, ETL pipelines, data modeling for legal and workflow systems.
- **Project Management:** Agile Methodologies, Sprint Planning, Task Tracking, Cross-functional Collaboration.